

1. data preprocessing

Handling miss value

```
import pandas as pd

import numpy as np

# Step 1: Load Dataset (Replace 'data.csv' with your actual file)
df = pd.read_csv("data.csv")

# Step 2: Display missing values count
print("Missing Values Before Handling:\n", df.isnull().sum())

# Step 3: Handling Missing Values

# 3.1 Fill missing values with Mean (for numerical columns)
df["Column1"] = df["Column1"].fillna(df["Column1"].mean())

print("\nMissing Values After Handling:\n", df.isnull().sum())
```

label encoding

```
label_encoder = LabelEncoder()

df["Category_LabelEncoded"] = label_encoder.fit_transform(df["Category"])

print("\nDataset after Label Encoding:\n", df)
```

Normalization

```
# Step 2: Apply Normalization (Min-Max Scaling)

minmax_scaler = MinMaxScaler()

df_normalized = pd.DataFrame(minmax_scaler.fit_transform(df), columns=df.columns)
```

2) logistic regression

```
import pandas as pd

import numpy as np

from sklearn.preprocessing

nd Living Area (GrLivArea)")

plt.ylabel("Sale Price")
```

```
plt.title("Simple Linear Regression")

# Residual Plot for Multiple Linear Regression
plt.subplot(1, 2, 2)
sns.residplot(x=y_pred_multiple, y=(y_test - y_pred_multiple), lowess=True, color="green")
plt.xlabel("Predicted Price")
plt.ylabel("Residuals")
plt.title("Multiple Linear Regression Residual Plot")
plt.tight_layout()
plt.show()

3) X = df[['GrLivArea', 'TotalBsmtSF', 'GarageArea']] # Features
y = df['PriceCategory'] # Binary Target (0 = Low Price, 1 = High Price)

# Step 4: Handle Missing Values
X.fillna(X.mean(), inplace=True)

# Step 5: Split Data into Training & Testing Sets (80% Train, 20% Test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 6: Standardize Features (for better performance)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Step 7: Train Logistic Regression Model
log_reg = LogisticRegression()
log_reg.fit(X_train, y_train)
y_pred = log_reg.predict(X_test)
y_prob = log_reg.predict_proba(X_test)[:, 1] # Probabilities for class 1
```

Step 8: Model Evaluation - Accuracy & Confusion Matrix

```
accuracy = accuracy_score(y_test, y_pred)
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
print(f"Model Accuracy: {accuracy:.2f}")
```

```
print("Confusion Matrix:\n", conf_matrix)
```

Step 9: ROC Curve & AUC Score

```
fpr, tpr, _ = roc_curve(y_test, y_prob)
```

```
roc_auc = auc(fpr, tpr)
```

Step 10: Plot ROC Curve

```
plt.figure(figsize=(8, 6))
```

```
plt.plot(fpr, tpr, color="blue", label=f"ROC Curve (AUC = {roc_auc:.2f})")
```

```
plt.plot([0, 1], [0, 1], color="grey", linestyle="--") # Diagonal line
```

```
plt.xlabel("False Positive Rate")
```

```
plt.ylabel("True Positive Rate")
```

```
plt.title("Receiver Operating Characteristic (ROC) Curve")
```

```
plt.legend(loc="lower right")
```

```
plt.show()
```

4)

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split, cross_val_score
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import accuracy_score, confusion_matrix


# Step 1: Load Dataset

file_path = "/mnt/data/House-Price.xlsx" # Update the path if needed
df = pd.read_excel(file_path)


# Step 2: Convert Problem into Binary Classification

# Predict whether SalePrice is high (above median) or low (below median)
df['PriceCategory'] = (df['SalePrice'] > df['SalePrice'].median()).astype(int)


# Step 3: Select Features & Target Variable

X = df[['GrLivArea', 'TotalBsmtSF', 'GarageArea']] # Features
y = df['PriceCategory'] # Binary Target (0 = Low Price, 1 = High Price)


# Step 4: Handle Missing Values

X.fillna(X.mean(), inplace=True)


# Step 5: Split Data into Training & Testing Sets (80% Train, 20% Test)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Step 6: Standardize Features (Important for KNN)

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)


# Step 7: Find Best k using Cross-Validation

k_values = range(1, 21)
```

```
cv_scores = []
```

```
for k in k_values:
```

```
    knn = KNeighborsClassifier(n_neighbors=k)
```

```
    scores = cross_val_score(knn, X_train, y_train, cv=5, scoring='accuracy')
```

```
    cv_scores.append(scores.mean())
```

```
# Best k value
```

```
best_k = k_values[np.argmax(cv_scores)]
```

```
print(f"Best k value: {best_k}")
```

```
# Step 8: Train KNN Model with Best k
```

```
knn_best = KNeighborsClassifier(n_neighbors=best_k)
```

```
knn_best.fit(X_train, y_train)
```

```
y_pred = knn_best.predict(X_test)
```

```
# Step 9: Model Evaluation
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
print(f"Model Accuracy: {accuracy:.2f}")
```

```
print("Confusion Matrix:\n", conf_matrix)
```

```
# Step 10: Plot k vs Accuracy
```

```
plt.figure(figsize=(8, 5))
```

```
plt.plot(k_values, cv_scores, marker='o', linestyle='dashed', color='blue', label="Cross-Validation Accuracy")
```

```
plt.xlabel("Number of Neighbors (k)")
```

```
plt.ylabel("Accuracy")
plt.title("Optimizing k in KNN Classification")
plt.legend()
plt.show()
```

```
5) import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
```

```
# Step 1: Load Dataset
```

```
file_path = "/mnt/data/House-Price.xlsx" # Update the path if needed
df = pd.read_excel(file_path)
```

```
# Step 2: Select Features for Clustering
```

```
X = df[['GrLivArea', 'TotalBsmtSF']] # Using two features for 2D visualization
```

```
# Step 3: Handle Missing Values
```

```
X.fillna(X.mean(), inplace=True)
```

```
# Step 4: Standardize Features
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

```
# Step 5: Find Optimal k using Elbow Method
```

```
wcss = [] # Within-cluster sum of squares
```

```
k_values = range(1, 11)
```

```
for k in k_values:
```

```
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
```

```
    kmeans.fit(X_scaled)
```

```
    wcss.append(kmeans.inertia_)
```

```
# Plot Elbow Method
```

```
plt.figure(figsize=(8, 5))
```

```
plt.plot(k_values, wcss, marker='o', linestyle='dashed', color='blue')
```

```
plt.xlabel("Number of Clusters (k)")
```

```
plt.ylabel("WCSS (Within-Cluster Sum of Squares)")
```

```
plt.title("Elbow Method for Optimal k")
```

```
plt.show()
```

```
# Step 6: Apply K-Means Clustering with Optimal k (Assume k=3 from elbow method)
```

```
optimal_k = 3
```

```
kmeans = KMeans(n_clusters=optimal_k, random_state=42, n_init=10)
```

```
df['Cluster'] = kmeans.fit_predict(X_scaled)
```

```
# Step 7: Visualize Clusters in 2D
```

```
plt.figure(figsize=(8, 6))
```

```
sns.scatterplot(x=X_scaled[:, 0], y=X_scaled[:, 1], hue=df['Cluster'], palette="viridis", s=100)
```

```
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], c='red', marker='X',  
s=200, label='Centroids')
```

```
plt.xlabel("GrLivArea (Standardized)")
```

```
pl
```

6)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, export_graphviz
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix
import graphviz
from IPython.display import display

# Step 1: Load Dataset
file_path = "/mnt/data/House-Price.xlsx" # Update the path if needed
df = pd.read_excel(file_path)

# Step 2: Convert Problem into Binary Classification
df['PriceCategory'] = (df['SalePrice'] > df['SalePrice'].median()).astype(int)

# Step 3: Select Features & Target Variable
X = df[['GrLivArea', 'TotalBsmtSF', 'GarageArea']] # Features
y = df['PriceCategory'] # Binary Target (0 = Low Price, 1 = High Price)

# Step 4: Handle Missing Values
X.fillna(X.mean(), inplace=True)

# Step 5: Split Data into Training & Testing Sets (80% Train, 20% Test)
```



```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Step 6: Train Decision Tree Model
```

```
dtree = DecisionTreeClassifier(max_depth=3, random_state=42) # Limit depth for better visualization
```

```
dtree.fit(X_train, y_train)
```

```
y_pred_tree = dtree.predict(X_test)
```

```
# Step 7: Train Random Forest Model
```

```
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
rf_model.fit(X_train, y_train)
```

```
y_pred_rf = rf_model.predict(X_test)
```

```
# Step 8: Model Evaluation
```

```
accuracy_tree = accuracy_score(y_test, y_pred_tree)
```

```
accuracy_rf = accuracy_score(y_test, y_pred_rf)
```

```
print(f"Decision Tree Accuracy: {accuracy_tree:.2f}")
```

```
print(f"Random Forest Accuracy: {accuracy_rf:.2f}")
```

```
# Step 9: Visualize Decision Tree using Graphviz
```

```
dot_data = export_graphviz(dtree, out_file=None, feature_names=X.columns,  
class_names=["Low Price", "High Price"],
```

```
filled=True, rounded=True, special_characters=True)
```

```
graph = graphviz.Source(dot_data)
```

```
display(graph) # Displays the decision tree visualization
```

```
7)
```

```
import pandas as pd
```

```
import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.preprocessing import StandardScaler

from sklearn.metrics import accuracy_score, confusion_matrix


# Step 1: Load Dataset

file_path = "/mnt/data/House-Price.xlsx" # Update the path if needed

df = pd.read_excel(file_path)


# Step 2: Convert Problem into Binary Classification

df['PriceCategory'] = (df['SalePrice'] > df['SalePrice'].median()).astype(int)


# Step 3: Select Features & Target Variable

X = df[['GrLivArea', 'TotalBsmtSF']] # Using 2 features for visualization

y = df['PriceCategory'] # Binary Target (0 = Low Price, 1 = High Price)


# Step 4: Handle Missing Values

X.fillna(X.mean(), inplace=True)


# Step 5: Split Data into Training & Testing Sets (80% Train, 20% Test)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Step 6: Standardize Features (Important for SVM)

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

```
# Step 7: Train SVM with Linear Kernel
```

```
svm_linear = SVC(kernel='linear', random_state=42)
```

```
svm_linear.fit(X_train, y_train)
```

```
y_pred_linear = svm_linear.predict(X_test)
```

```
# Step 8: Train SVM with RBF Kernel
```

```
svm_rbf = SVC(kernel='rbf', gamma='scale', random_state=42)
```

```
svm_rbf.fit(X_train, y_train)
```

```
y_pred_rbf = svm_rbf.predict(X_test)
```

```
# Step 9: Model Evaluation
```

```
accuracy_linear = accuracy_score(y_test, y_pred_linear)
```

```
accuracy_rbf = accuracy_score(y_test, y_pred_rbf)
```

```
print(f"SVM Linear Kernel Accuracy: {accuracy_linear:.2f}")
```

```
print(f"SVM RBF Kernel Accuracy: {accuracy_rbf:.2f}")
```

```
# Step 10: Plot Decision Boundaries
```

```
def plot_decision_boundary(model, X, y, title):
```

```
    h = 0.02 # Step size
```

```
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
```

```
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
```

```
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
```

```
                        np.arange(y_min, y_max, h))
```

```
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
```

```
    Z = Z.reshape(xx.shape)
```

```
plt.figure(figsize=(8, 6))  
plt.contourf(xx, yy, Z, alpha=0.3, cmap="coolwarm")  
sns.scatterplot(x=X[:, 0], y=X[:, 1], hue=y, palette="coolwarm", edgecolor="black")  
plt.xlabel("GrLivArea (Standardized)")  
plt.ylabel("TotalBsmtSF (Standardized)")  
plt.title(title)  
plt.show()
```

Plot Decision Boundaries for both models

```
plot_decision_boundary(svm_linear, X_train, y_train, "SVM Linear Kernel Decision  
Boundary")
```

```
plot_decision_boundary(svm_rbf, X_train, y_train, "SVM RBF Kernel Decision Boundary")
```

8)

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.decomposition import PCA
```

```
from sklearn.preprocessing import StandardScaler
```

```
from mpl_toolkits.mplot3d import Axes3D
```

Step 1: Load Dataset

```
file_path = "/mnt/data/House-Price.xlsx" # Update path if needed
```

```
df = pd.read_excel(file_path)
```

Step 2: Select Numerical Features for PCA

```
features = ['GrLivArea', 'TotalBsmtSF', 'GarageArea', 'LotArea', 'YearBuilt']
```

```
X = df[features]
```

```
# Step 3: Handle Missing Values
```

```
X.fillna(X.mean(), inplace=True)
```

```
# Step 4: Standardize Data
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

```
# Step 5: Apply PCA to Reduce to 3 Components
```

```
pca = PCA(n_components=3)
```

```
X_pca = pca.fit_transform(X_scaled)
```

```
# Step 6: Convert PCA Output to DataFrame
```

```
df_pca = pd.DataFrame(X_pca, columns=['PC1', 'PC2', 'PC3'])
```

```
# Step 7: 3D Visualization
```

```
fig = plt.figure(figsize=(10, 7))
```

```
ax = fig.add_subplot(111, projection='3d')
```

```
ax.scatter(df_pca['PC1'], df_pca['PC2'], df_pca['PC3'], c='blue', marker='o', alpha=0.6)
```

```
# Labels & Title
```

```
ax.set_xlabel('Principal Component 1')
```

```
ax.set_ylabel('Principal Component 2')
```

```
ax.set_zlabel('Principal Component 3')
```

```
ax.set_title('3D Visualization of PCA-transformed Data')
```

```
plt.show()
```

```
# Step 8: Explained Variance Ratio
```

```
print("Explained Variance Ratio:", pca.explained_variance_ratio_)
```

```
9)
```

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
import re
```

```
import string
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
from sklearn.naive_bayes import MultinomialNB
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
# Step 1: Load Sample Sentiment Dataset
```

```
data = {
```

```
    "text": [
```

```
        "I love this product! It's amazing.",
```

```
        "This is the worst experience I've ever had.",
```

```
        "Absolutely fantastic! I would buy it again.",
```

```
        "I hate this so much. Waste of money!",
```

```
        "Not bad, but could be better.",
```

```
        "The quality is terrible. Never again!",
```

```
        "I'm very happy with my purchase.",
```

```
        "Awful experience, totally disappointed.",
```

```
"Decent product for the price.",  
"Horrible! I regret buying this."  
],  
"sentiment": [1, 0, 1, 0, 1, 0, 1, 0, 1, 0] # 1 = Positive, 0 = Negative  
}
```

```
df = pd.DataFrame(data)
```

```
# Step 2: Text Preprocessing Function
```

```
def clean_text(text):
```

```
    text = text.lower() # Lowercase
```

```
    text = re.sub(f"[{string.punctuation}]", "", text) # Remove punctuation
```

```
    return text
```

```
df["clean_text"] = df["text"].apply(clean_text)
```

```
# Step 3: Convert Text to Numerical Features (TF-IDF Vectorization)
```

```
vectorizer = TfidfVectorizer()
```

```
X = vectorizer.fit_transform(df["clean_text"])
```

```
y = df["sentiment"]
```

```
# Step 4: Split Data into Training & Testing Sets (80% Train, 20% Test)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Step 5: Train Naïve Bayes Model
```

```
nb_classifier = MultinomialNB()
```

```
nb_classifier.fit(X_train, y_train)
```

```
# Step 6: Make Predictions
```

```
y_pred = nb_classifier.predict(X_test)
```

```
# Step 7: Model Evaluation
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
report = classification_report(y_test, y_pred)
```

```
print(f"Naïve Bayes Classifier Accuracy: {accuracy:.2f}")
```

```
print("\nConfusion Matrix:\n", conf_matrix)
```

```
print("\nClassification Report:\n", report)
```

```
10)
```

```
import tensorflow as tf
```

```
from tensorflow import keras
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
# Step 1: Load MNIST Dataset
```

```
(X_train, y_train), (X_test, y_test) = keras.datasets.mnist.load_data()
```

```
# Step 2: Normalize Pixel Values (Scale between 0 and 1)
```

```
X_train, X_test = X_train / 255.0, X_test / 255.0
```

```
# Step 3: Build ANN Model
```

```
model = keras.Sequential([
```

```
    keras.layers.Flatten(input_shape=(28, 28)), # Flatten 28x28 images to 1D
```



```
keras.layers.Dense(128, activation='relu'), # Hidden Layer (128 neurons, ReLU activation)

keras.layers.Dense(10, activation='softmax') # Output Layer (10 neurons for digits 0-9,
Softmax activation)

])
```

Step 4: Compile Model

```
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
metrics=['accuracy'])
```

Step 5: Train the Model

```
history = model.fit(X_train, y_train, epochs=10, validation_data=(X_test, y_test))
```

Step 6: Evaluate Model Performance

```
test_loss, test_acc = model.evaluate(X_test, y_test, verbose=2)
```

```
print(f"\nTest Accuracy: {test_acc:.2f}")
```

Step 7: Predictions

```
y_pred = np.argmax(model.predict(X_test), axis=1)
```

Step 8: Classification Report

```
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Step 9: Confusion Matrix

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
plt.figure(figsize=(10, 8))
```

```
plt.imshow(conf_matrix, cmap='Blues', interpolation='nearest')
```

```
plt.colorbar()
```

```
plt.title('Confusion Matrix')
```

```
plt.xlabel('Predicted Label')
```

```
plt.ylabel('True Label')
```

```
plt.show()
```

```
# Step 10: Visualize Some Predictions
```

```
def plot_images(X, y_true, y_pred, num_images=10):
```

```
    plt.figure(figsize=(10, 5))
```

```
    for i in range(num_images):
```

```
        plt.subplot(2, 5, i + 1)
```

```
        plt.imshow(X[i], cmap='gray')
```

```
        plt.title(f"True: {y_true[i]}\nPred: {y_pred[i]}")
```

```
        plt.axis("off")
```

```
    plt.show()
```

```
plot_images(X_test[:10], y_test[:10], y_pred[:10])
```

```
11)
```

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.ensemble import IsolationForest
```

```
from sklearn.svm import OneClassSVM
```

```
from sklearn.preprocessing import StandardScaler
```

```
# Step 1: Load Sample Data (Simulating Normal & Anomalous Data)
```

```
np.random.seed(42)
```

```
normal_data = np.random.normal(loc=50, scale=10, size=(200, 2)) # Normal Data
```

```
anomalous_data = np.random.normal(loc=80, scale=5, size=(10, 2)) # Anomalous Data
data = np.vstack((normal_data, anomalous_data))

# Convert to DataFrame
df = pd.DataFrame(data, columns=["Feature1", "Feature2"])

# Step 2: Normalize Data
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df)

# Step 3: Train Isolation Forest Model
iso_forest = IsolationForest(n_estimators=100, contamination=0.05, random_state=42)
iso_preds = iso_forest.fit_predict(df_scaled)

# Step 4: Train One-Class SVM Model
oc_svm = OneClassSVM(kernel="rbf", nu=0.05)
oc_preds = oc_svm.fit_predict(df_scaled)

# Step 5: Convert Predictions (-1: Anomaly, 1: Normal)
df["IsoForest_Label"] = iso_preds
df["OneClassSVM_Label"] = oc_preds

# Step 6: Visualize Anomalies Detected by Isolation Forest
plt.figure(figsize=(10, 6))

sns.scatterplot(x=df["Feature1"], y=df["Feature2"], hue=df["IsoForest_Label"], palette={1: 'blue', -1: 'red'})

plt.title("Anomaly Detection using Isolation Forest")
plt.legend(["Normal", "Anomaly"])
```

```
plt.show()
```

```
# Step 7: Visualize Anomalies Detected by One-Class SVM
```

```
plt.figure(figsize=(10, 6))
```

```
sns.scatterplot(x=df["Feature1"], y=df["Feature2"], hue=df["OneClassSVM_Label"],  
palette={1: 'blue', -1: 'red'})
```

```
plt.title("Anomaly Detection using One-Class SVM")
```

```
plt.legend(["Normal", "Anomaly"])
```

```
plt.show()
```